



Machine Learning for Time Series

(MLTS or MLTS-Deluxe Lectures)

Dr. Dario Zanca

Machine Learning and Data Analytics (MaD) Lab
Friedrich-Alexander-Universität Erlangen-Nürnberg
25.10.2022


-
- Time series fundamentals and definitions (2 lectures)
 - Bayesian Inference (1 lecture)
 - Gaussian processes (2 lectures)
 - State space models (2 lectures)
 - Autoregressive models (1 lecture)
 - Data mining on time series (1 lecture)
 - Deep learning on time series (4 lectures) ←
 - Domain adaptation (1 lecture)
-

Recap: Recurrent neural networks

RNN / LSTM limitations:

- Non-parallelism → Long training time
- Difficulties with long sequences
 - Large memory usage
 - Difficult to train (vanishing/exploding gradients)
 - Hard to learn long-term dependencies (mitigated by LSTMs)

Motivation



- Transformers *perceive* the entire sequence at the same time.
- They are based on the seq2seq concept, i.e., transforming sequences into other sequences.
- State of the art in many NLP tasks.

In this lecture...

- Attention models
 - The transformers architecture
-



Deep Learning for Time Series – Attention models

Attention models



Sequence-to-sequence models

Sequence-to-sequence (seq2seq) models aim at transforming an input sequence to an output sequence

- E.g., machine translation.

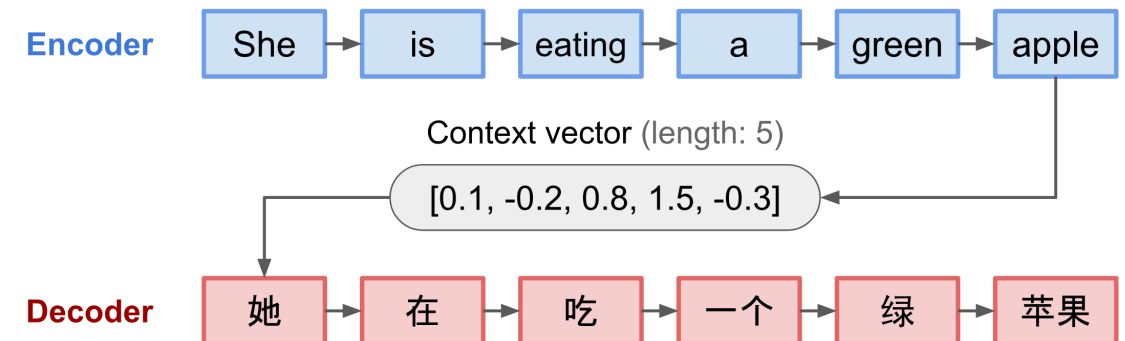
Seq2seq models generally have an encoder-decoder architecture, composed by:

- An encoder that processes the input sequence and compress the information into a context vector (also said embedding).
 - A decoder that processes the context vector and produces the transformed output.
-

Sequence-to-sequence models

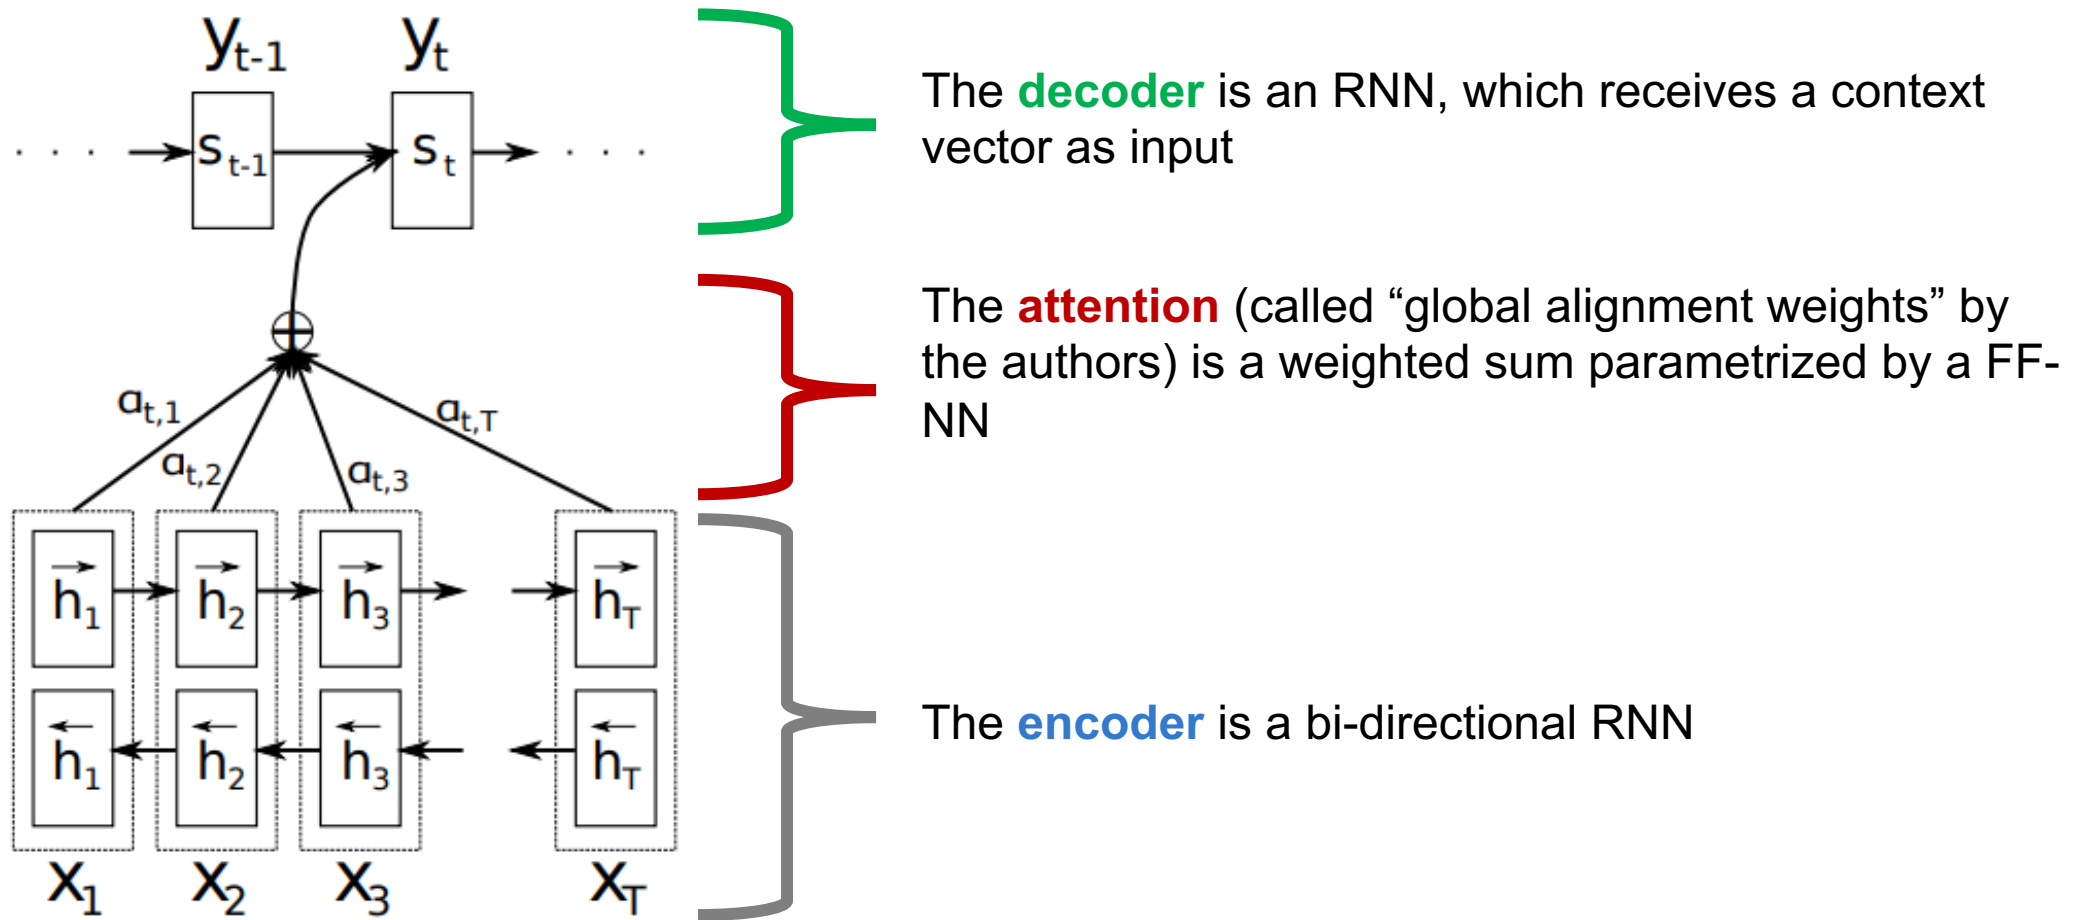
Disadvantages of the fixed length context vector design are:

- Incapability of remembering long sequences
 - Too complex dynamics to be encoded in the hidden state
- Attention mechanisms are proposed to solve this problem.
- Originated for machine translation [1]



[1] "Neural machine translation by jointly learning to align and translate", Bahdanau et al.
Image from: <https://lilianweng.github.io/lil-log/2018/06/24/attention-attention.html>

Attention mechanism



Attention mechanism: formalization

Let x be the input sequence of length n , and y the output sequence of length m .

Let $h_i = [\overrightarrow{h_i}, \overleftarrow{h_i}]$ be the encoder state, given by the concatenation of the forward and backward hidden states of the bidirectional RNN.

Let denote with s_t the decoder state, for the output at position t . Then, the context vector is defined as the sum of the encoder states, weighted by the alignment scores, i.e.,


$$c_t = \sum_{i=1}^n a_{t,i} h_i$$

where $a_{t,i} = \text{softmax}(\text{score}(s_{t-1}, h_i))$, and $\text{score}(s_{t-1}, h_i) = v_a \tanh(W_a[s_{t-1}; h_i])$.

Attention mechanism: example

In the example on the side, the x-axis represents the source sentence and the y-axis the generated sentence.

Every pixel (i, j) , at the i -th row and j -th column, indicates the weight the j -th source word embedding when generating the i -th word.



Attention mechanism

Name	Alignment score function
Content-base attention	$\text{score}(s_t, h_i) = \text{cosine}[s_t, h_i]$
Additive(*)	$\text{score}(s_t, h_i) = \mathbf{v}_a^\top \tanh(\mathbf{W}_a[s_t; h_i])$
Location-Base	$\alpha_{t,i} = \text{softmax}(\mathbf{W}_a s_t)$ Note: This simplifies the softmax alignment to only depend on the target position.
General	$\text{score}(s_t, h_i) = s_t^\top \mathbf{W}_a h_i$ where \mathbf{W}_a is a trainable weight matrix in the attention layer.
Dot-Product	$\text{score}(s_t, h_i) = s_t^\top h_i$
Scaled Dot-Product(^)	$\text{score}(s_t, h_i) = \frac{s_t^\top h_i}{\sqrt{n}}$ Note: very similar to the dot-product attention except for a scaling factor; where n is the dimension of the source hidden state.

Table from: <https://lilianweng.github.io/lil-log/2018/06/24/attention-attention.htm>

Attention mechanism

Attention mechanisms can be characterised according to their function or design as:

- **Self-attention**
 - **Soft/Hard attention**
 - **Global/Local attention**
-


Global / local attention

Self-attention, also said intra-attention, is an attention mechanisms which relates different positions of an input sequence to generate a representation of the same sequence.

The FBI is chasing a criminal on the run .
The FBI is chasing a criminal on the run .
The FBI is chasing a criminal on the run .
The FBI is chasing a criminal on the run .
The FBI is chasing a criminal on the run .
The FBI is chasing a criminal on the run .
The FBI is chasing a criminal on the run .
The FBI is chasing a criminal on the run .
The FBI is chasing a criminal on the run .
The FBI is chasing a criminal on the run .


Soft / hard attention

Soft attention applies learned weights to input parts, while **hard attention** select a single input part at the time.



Global / local attention

In **global attention**, the context vector depends on the whole input sequence. The **local attention**, generates an input vector which depends only on a subset of the input sequence, corresponding to a window centered on the current position.





Deep Learning for Time Series – Attention models

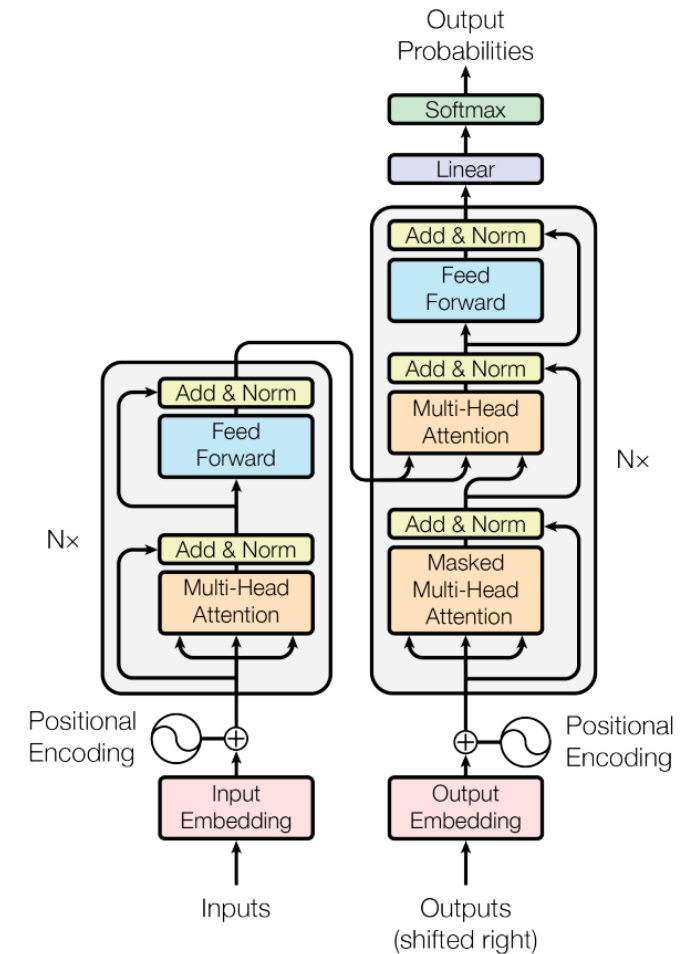
The Transformer architecture



The Transformer architecture

The Transformer architecture, introduced in 2017 [2]:


- Completely built on self-attention
 - Do not use sequence aligned recurrent architecture
- Replaced all the RNNs





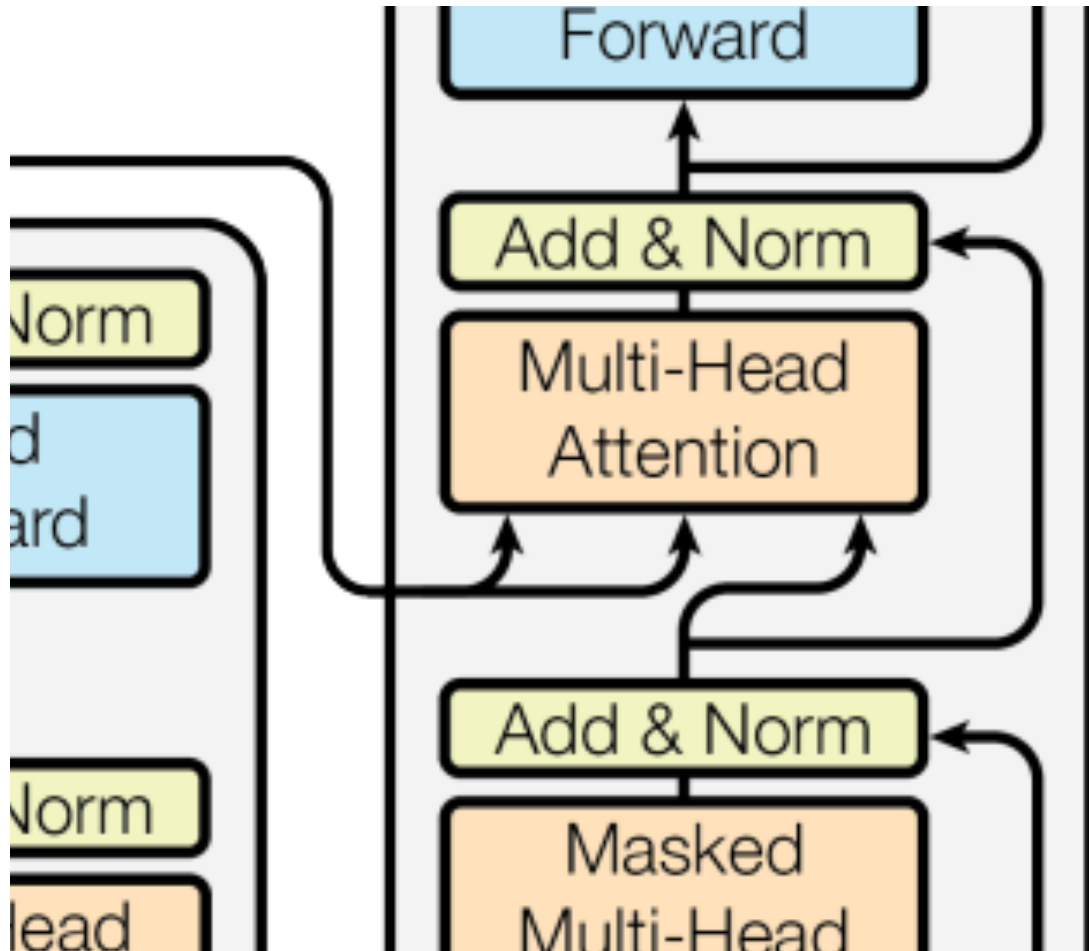
The Transformer architecture

The fundamental components of the transformer architectures are:

- Positional encoding
- Multi-head self attention
 - Based on K, V, and Q matrices
- An encoder-decoder architecture




Transformer: multi-head attention



Multi-head attention

Transformer: multi-head self-attention



Multi-head attention

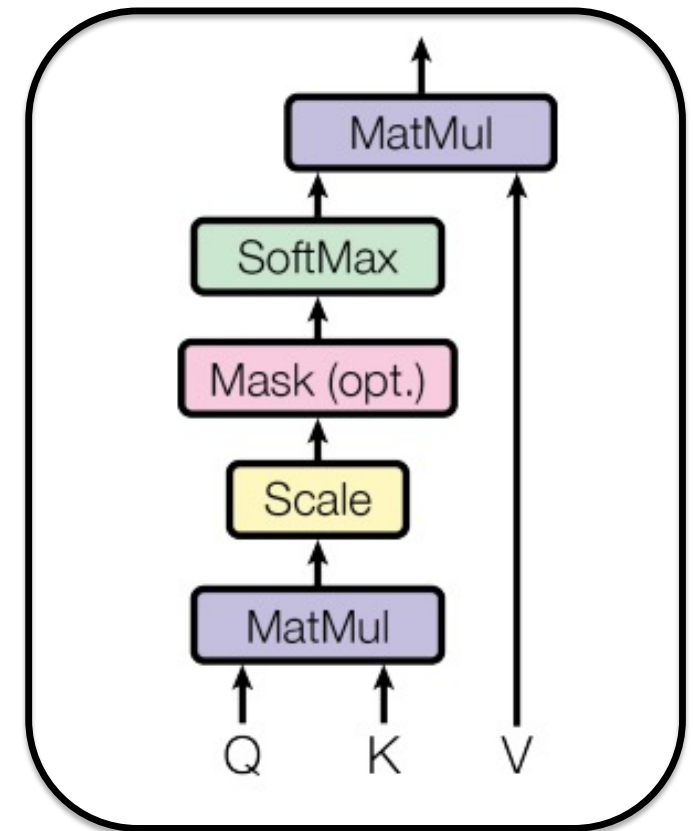
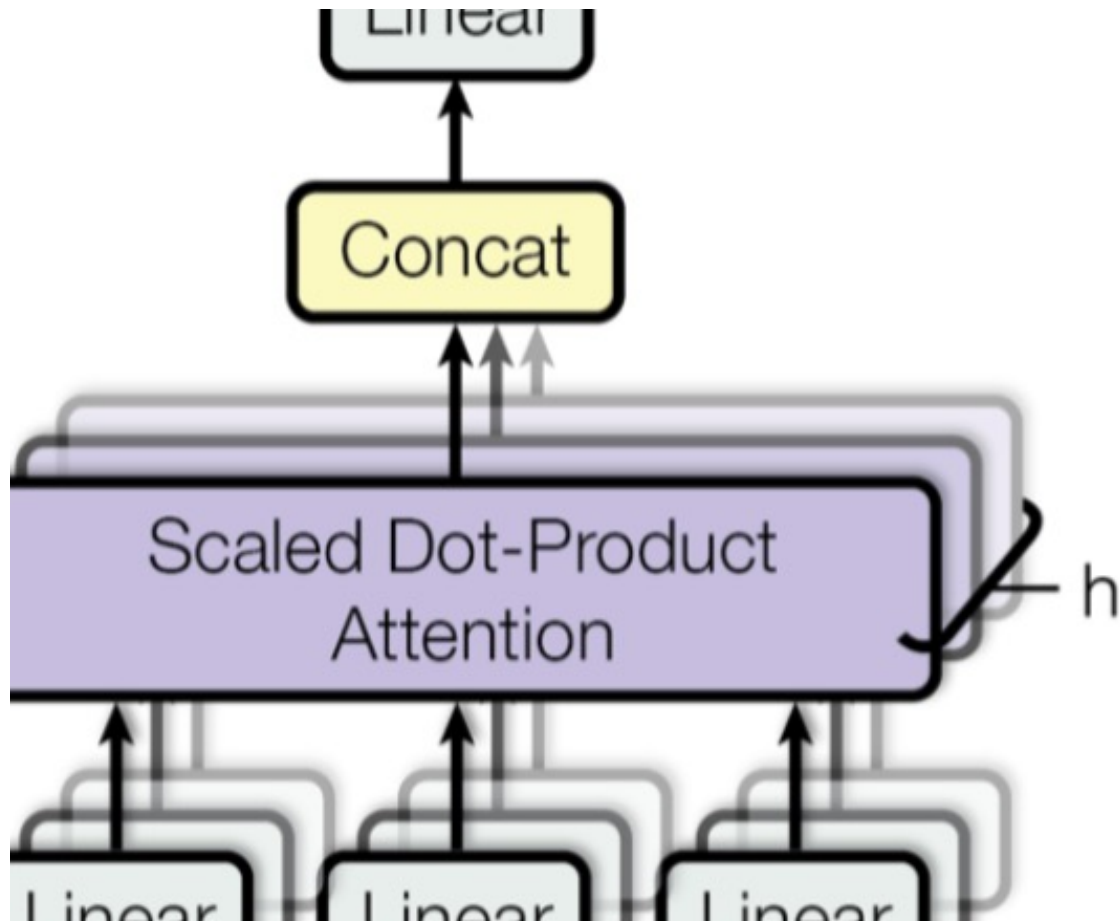
[2] "Attention is all you need", Vaswani, et al.

The first step in calculating self-attention consists of calculating three vectors from each input vector (i.e., each element of the input sequence).

- Query, Q
- Key, K
- Value, V


These vector are calculated by multiplying the input vector by a corresponding matrix, resp., W^Q , W^K , and W^V .

Transformer: multi-head self-attention



Scaled Dot-Product Attention

Transformer: multi-head self-attention




Scaled Dot-Product Attention

Then, the scaled dot-product attention is computed by

$$Attention(Q, K, V) = softmax\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

where d_k is the dimension of queries and keys.

Transformer: multi-head self-attention



Multi-head attention

[2] “Attention is all you need”, Vaswani, et al.

Instead of performing a single attention function with a d_{model} -dimensional keys, values and queries, it is beneficial to linearly project queries, keys and values with **different linear projections**.

Each of this projections is processed **in parallel** and then **concatenated**.

$$MultiHead(Q, K, V) = concat(head_1, \dots, head_h)$$

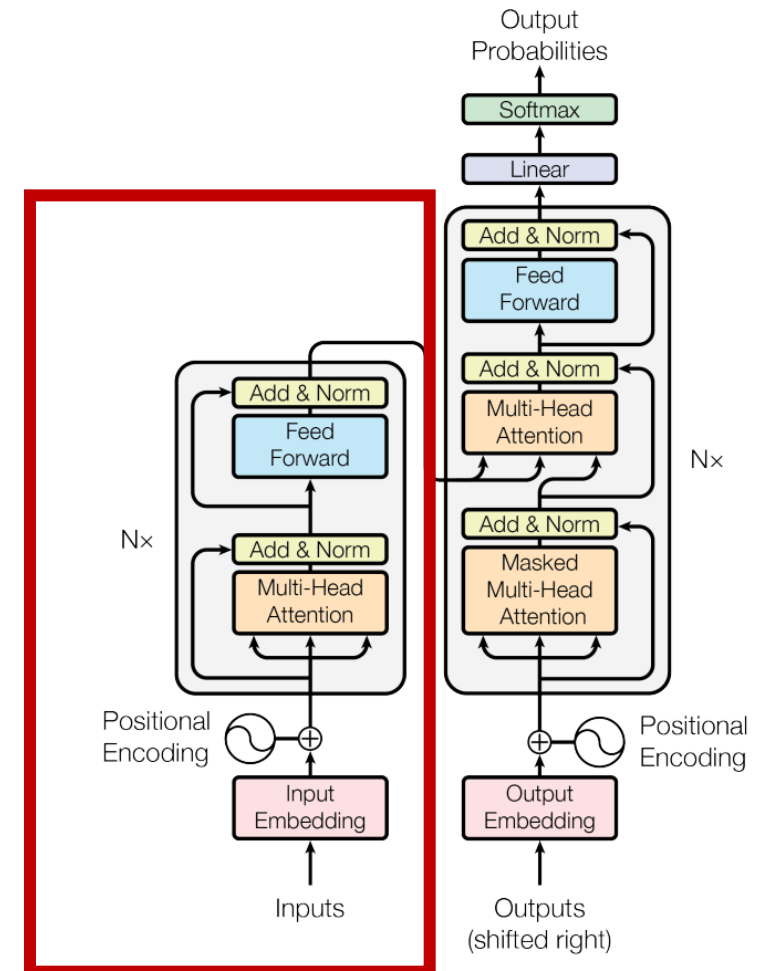
where each of the heads is a scaled dot-product attention, i.e.,

$$head_i = Attention(QW_i^Q, KW_i^K, VW_i^V)$$

Transformer: the encoder

The encoder generates an attention-based representation of the input sequence.

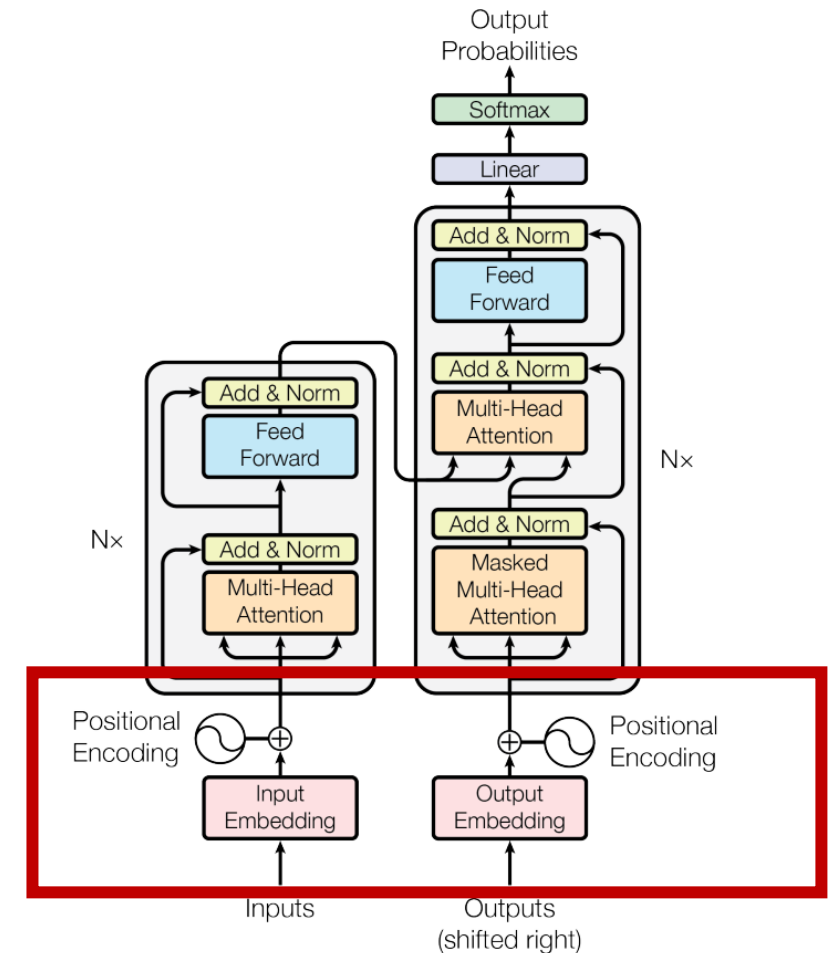
- Capability of locate a specific information from a (potentially) very large context.
- It is made of N_x identical layers,, each composed of
 - A multi-head attention layer
 - A positional FF-NN
 - A residual connection and layer normalization



Transformer: the positional encoding

Since the model contains no recurrence nor convolution, in order to make use of the order of the sequence, positional information is injected by mean of a **positional encoding**.

The positional encoding is **added to the input vectors** (for both the encoder and the decoder networks).



Transformer: the positional encoding


In the original implementations, the positional encoding is defined by using sine and cosine functions of different frequencies:

$$PE_{(pos, 2i)} = \sin\left(\frac{pos}{10000^{\frac{2i}{d_{model}}}}\right)$$


$$PE_{(pos, 2i+1)} = \cos\left(\frac{pos}{10000^{\frac{2i}{d_{model}}}}\right)$$

where pos is the position and i is the dimension.

[2] “Attention is all you need”, Vaswani, et al.




Transformer: the positional encoding



Code to generate this example available at:
https://github.com/jalammar/jalammar.github.io/blob/master/notebooks/transformer/transformer_positional_encoding_graph.ipynb


Transformer: the positional encoding



Positional encoding of the 1st input

Code to generate this example available at:
https://github.com/jalammar/jalammar.github.io/blob/master/notebooks/transformer/transformer_positional_encoding_graph.ipynb

Transformer: the positional encoding



Code to generate this example available at:
https://github.com/jalammar/jalammar.github.io/blob/master/notebooks/transformer/transformer_positional_encoding_graph.ipynb

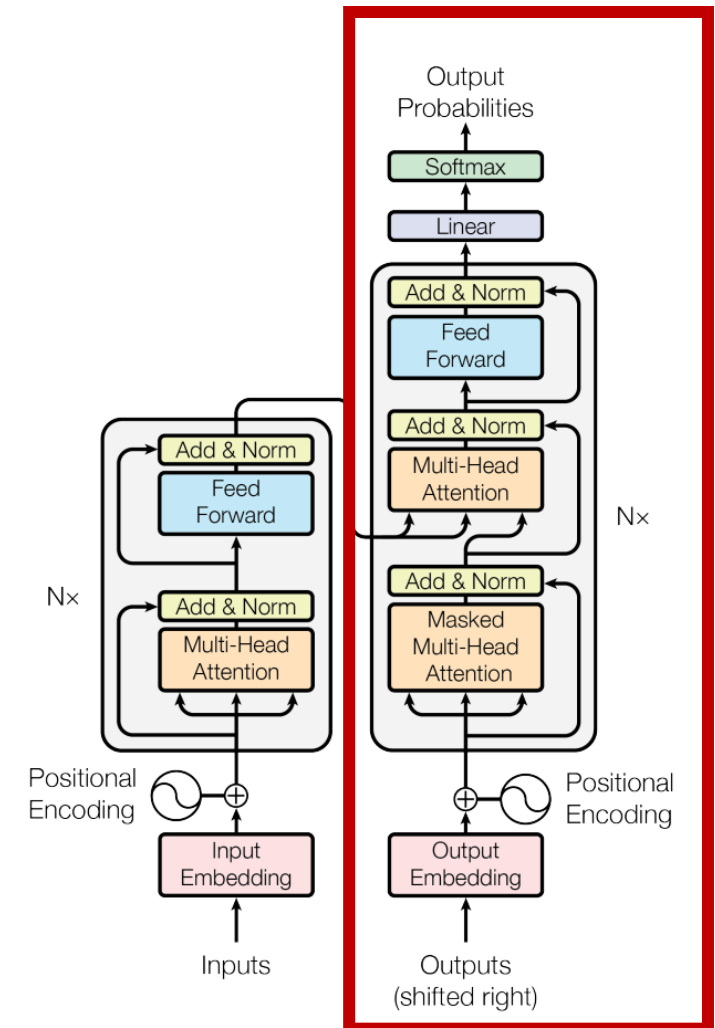
Positional encoding of the 10th input

Transformer: the decoder

The decoder network is able to retrieval information from the encoded representation.

- As the encoder, it is made of N_x identical layers,, each composed of
 - A multi-head attention layer
 - A positional FF-NN
 - A residual connection and layer normalization
- The first multi-head self-attention is properly masked to prevent information leakage from future positions.


[2] “Attention is all you need”, Vaswani, et al.



Transformer: the decoder

The decoding works as follows:

1. The output of the top encoder is transformed into a set of attention vectors K and V .
2. These are used by each decoder
→ It helps the decoder focus on appropriate parts
3. The process is repeated until a special symbol of <end> is generated.





Lecture title

Recap



In this lecture

- Attention models
 - Seq2Seq
 - Encoder-decoder
 - All hidden states are passed
 - Self-attention, soft/hard, local/global
 - Transformers
 - Multi-head Self-attention
 - Positional encoding
 - Encoder and decoder networks
-

Transformers: pros and cons

Pros:

- Transformers can learn direct access to potentially very far input parts
- Many degree of freedom → Can learn complex dependencies
- Feed-forward model provides high performance on modern hardware

Cons:

- Quadratic time and memory complexity
- Many degree of freedom → Data hungry behavior during training
- Training is insidious
 - Hard integration with other architecture, not easy learning rate policy, use custom data loader, ...
- Still limited research on time series data.

Other attention-based approaches

- Neural Turing Machines [3]
 - Couples NNs with external memory
 - Mimics reading and writing operations in Turing machines
 - Exploit content-based attention
 - Pointer networks [4]
 - Solves problems where the outputs are positions in an input sequence
 - It applies attention over the input elements to pick one as the output at each decoder step.
 - ...
-

